

Predicting the Zero-Shot Transfer of a Promptable Video Tracker

Forecasting SAM 3's per-species, per-place reliability —
for detection and association — from label-free distances

Kostiantyn Boiar
Student ID: <TODO>

Submitted in partial fulfilment of the requirements for the
Degree of Master of Science in <TODO: programme>

School of Computing Science
University of Glasgow

Supervisor: Dr Marwa Mahmoud

July 28, 2026

Education Use Consent & Declaration

Education Use Consent.

I hereby give my permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format.

Name: _____ Signature: _____

Declaration of Originality.

I declare that this dissertation is my own work, except where explicitly stated and referenced, and that it has not been submitted for any other degree at the University of Glasgow or any other institution.

Abstract

Promptable trackers such as SAM3 segment and track any named animal in a video with no training on that species, yet no rule tells a practitioner *in advance* whether the model is reliable for a given animal in a given place. This dissertation asks whether such a rule exists. On the SA-FARI benchmark, whose splits are disjoint by species and location, we measure SAM3’s zero-shot promptable-tracking accuracy with the official evaluator, decompose it into detection (pDetA) and association (pAssA), and test whether each is predicted, *before running the model* and without any target label, by four distances (taxonomic, visual, environment, temporal). Under a deliberately powerful test — whole-species and whole-location hold-out with group-cluster-bootstrap significance — **it is not**: no distance beats a mean-predictor baseline out of sample, and the one apparent effect is an animal-size confound. A follow-up pivot to intrinsic scene difficulty (low-light, clutter) collapses to the same size effect, and the decomposition shows association carries almost no signal distinct from detection. The contribution is thus a rigorous, decomposed *negative* result — the first for a promptable video tracker — validated by a test shown to detect the one real signal (object size) when it is present.

Acknowledgements

I thank my supervisor, Dr Marwa Mahmoud, for guidance throughout this project.

Contents

Abstract	ii
Acknowledgements	iii
1 Introduction	1
1.1 Motivation	1
1.2 Problem statement and research question	1
1.3 Hypotheses	1
1.4 Contributions	1
1.5 Dissertation outline	2
2 Background and Related Work	3
2.1 Promptable segmentation and tracking	3
2.2 The tracking metric and its decomposition	3
2.3 Predictable out-of-distribution generalisation	3
2.4 Distribution shift in camera-trap and wildlife vision	4
2.5 What should predict transfer	4
2.6 The SA-FARI benchmark	4
2.7 Summary and gap	5
3 Data and Methodology	6
3.1 Research design and evaluation criteria	6
3.2 Dataset, splits, and the analysis unit	6
3.3 Measuring transfer	7
3.4 The frozen reference (leakage firewall)	8
3.5 Label-free distance features	8
3.5.1 Taxonomic distance	8
3.5.2 Visual distance	8
3.5.3 Environment distance	8
3.5.4 Temporal gap	8
3.5.5 Animal size (confound control)	8
3.5.6 SAM3 familiarity proxy (deferred)	9
3.6 Statistical modelling and hypothesis testing	9
3.6.1 Per-target regression	9
3.6.2 Honest, group-aware inference	9
3.6.3 Out-of-sample validation and the significance test	9
3.6.4 Variance attribution and confound control	10
3.6.5 The intrinsic-difficulty pivot and the nested decomposition	10
3.6.6 Design-robustness checks	10
3.7 Detection precision: false positives on hard negatives	10
3.8 Independent replication on external datasets	10
3.9 Reproducibility, implementation, and constraints	11

4	Results and Evaluation	12
4.1	Measurement	12
4.2	The fitted law	12
4.3	Variance partitioning and the intrinsic-difficulty decomposition	13
4.4	Out-of-sample validation	14
4.5	Species hold-out (H1): novelty and the size confound	15
4.6	False positives / hallucination	17
4.7	Ablations and robustness	17
4.8	Independent replication (MammAlps)	18
4.9	Independent replication (BURST)	19
4.10	Cross-dataset synthesis	21
5	Discussion	23
5.1	Interpretation	23
5.2	Relation to prior label-free evaluation	23
5.3	Why label-free distances do not yield a deployment estimator	24
5.4	Limitations and threats to validity	24
5.5	Ethical considerations	24
6	Conclusion	25
6.1	Summary of findings	25
6.2	Contributions	25
6.3	Future work	25
	Bibliography	26
A	Mathematical definitions	28
A.1	The promptable-tracking metric	28
A.2	Label-free distances	28
A.3	Model and statistics	29
B	Algorithms	30
C	Supplementary results	32

List of Figures

3.1	The measurement-and-modelling pipeline	7
3.2	The image-based distances in action	9
4.1	Coefficient forest plot	13
4.2	Out-of-sample predicted vs actual	15
4.3	The size confound	16

List of Tables

4.1	Zero-shot SAM3 on the SA-FARI test split — support-weighted mean over 346 scored cells (of 1447).	12
4.2	Standardised logit-GLM coefficients with <i>group-cluster-bootstrap</i> 95% CIs (resampling whole species / locations, so the interval respects pseudo-replication and the support weighting), for detection (pDetA) and association (pAssA). Every distance’s CI spans zero.	12
4.3	Variance partition for detection (pDetA): each distance’s unique share of R^2 (LMG/Shapley) and its VIF. The four distances together explain only $\approx 3\%$ of the variance, no factor dominates, and every $VIF \approx 1$ — the null is not masked by collinearity.	13
4.4	Nested size decomposition of the intrinsic-difficulty test (detection, 346 cells). Each entry is the out-of-sample MAE gain over a mean-predictor baseline, Δ [95% CI] (p), under leave-species-out and leave-location-out. The entire gain sits in the size covariate (log-area); low-light <i>alone</i> does not validate — the difficulty signals only pass the bar by riding on size.	14
4.5	Out-of-sample error: whole-group hold-out (leave-species-out / leave-location-out) vs a mean predictor. Δ is the MAE gain over baseline; its bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. No hold-out’s gain is significant.	14
4.6	Zero-shot SAM3 on the pooled train+test set — support-weighted mean over 1025 scored cells (of 2126).	15
4.7	Standardised logit-GLM coefficients with <i>group-cluster-bootstrap</i> 95% CIs (resampling whole species / locations, so the interval respects pseudo-replication and the support weighting), for detection (pDetA) and association (pAssA).	16
4.8	Out-of-sample error: whole-group hold-out (leave-species-out / leave-location-out) vs a mean predictor. Δ is the MAE gain over baseline; its bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. No hold-out’s gain is significant.	16
4.9	False positives on the 1486 hard negatives (113 species): the overall hallucination rate is 9.8%. More visually distinctive species hallucinate <i>less</i> — a real, size-independent correlate — but it does not validate out of sample (leave-species-out $p = 0.053$), so it is a correlate, not a predictor.	17
4.10	Factor ablation for detection (pDetA): McFadden pseudo- R^2 and the out-of-sample MAE gain over a mean baseline under leave-species-out (LSO) and leave-location-out (LLO). Dropping any single distance barely changes the fit or the (null) out-of-sample gain — no factor carries signal; only adding the animal-size covariate moves the score.	17

4.11	Design-robustness sweep (detection). Each row re-runs the primary four-distance model with one design choice swapped; entries are the out-of-sample MAE gain over a mean-predictor baseline, Δ [95% CI] (p), under leave-species-out and leave-location-out, with significance judged after Bonferroni ($m = 5$). The null holds throughout: no encoder, crop, prompt, or distributional distance makes the label-free distances beat the mean.	18
4.12	Zero-shot SAM3 on MammAlps — support-weighted mean over the 20 probe cells (MammAlps ships no separate reference split; all clips are probe). Scored with bbox -HOTA (<code>eval.prefer_bbox</code>): the ground truth is box-only, so mask-IoU is near-zero (dataset mask mAP ≈ 0.002) and box-IoU is the honest metric. The score is non-degenerate — SAM3 genuinely finds Alpine mammals zero-shot. . .	19
4.13	MammAlps independent replication — out-of-sample error of the pre-registered label-free distance model (taxonomic + visual + environment) under whole-group hold-out, against a mean predictor. Hold-out <i>camera</i> is leave-one-camera-out (the environment scheme); <i>species</i> is leave-one-species-out. Δ is the MAE gain over baseline; the bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. Every interval spans zero on both targets and both schemes: the label-free distances do not predict transfer, directionally consistent with the SA-FARI null. $n = 20$ cells (5 species, 9 cameras); a single-cell fox fold and near-constant taxonomic distance leave the species scheme underpowered. . .	19
4.14	Zero-shot SAM3 on the BURST val animal subset — support-weighted mean over the 132 probe cells (41 species). Mask-native ground truth, scored with mask-HOTA (dataset mask mAP 0.49, bbox mAP 0.51). The score is non-degenerate and the association target is genuine ($\text{corr}(\mathbf{pDetA}, \mathbf{pAssA}) = 0.76$, 71/132 multi-object cells).	20
4.15	BURST independent replication — out-of-sample error under leakage-free leave-species-out cross-validation (41 species, 132 mask-native cells), against a mean predictor. Δ is the MAE gain over baseline for pDetA ; the bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out species. Every before-running distance is null (CI spans zero); the after-running confidence controls fire, but through a within-species channel a species-constant distance cannot use (collapsing <code>conf_mean_score</code> to species means drops it to +0.004, $p = 0.21$), so they establish pipeline liveness, not power for a distance. BURST has no shared locations, so leave-species-out is the only valid scheme.	20
4.16	What was tested, the numerical out-of-sample result, and on which dataset. Result is the leave-species-out gain in MAE over a mean predictor (ΔMAE , with one-sided bootstrap p) unless noted; positive and significant means the feature predicts transfer. The before-running label-free distances are null on all three datasets; the only feature that validates is SAM3’s own <i>after-running</i> confidence (and only decisively on SA-FARI — on BURST it fires through a within-species channel a distance cannot use, and on MammAlps the design is too small for anything to validate). “Opposite direction” means the association runs the reverse of what novelty predicts — more-novel animals are detected <i>better</i> , not worse — so it cannot support the hypothesis.	22

Chapter 1

Introduction

1.1 Motivation

Promptable video trackers such as SAM3 [4] segment and track any named animal in a video with no training on that species, and conservation teams now deploy them on new species and regions constantly. Yet no rule tells a practitioner, *before running the model*, whether its output is trustworthy for a given animal in a given place.

1.2 Problem statement and research question

This dissertation asks whether SAM3’s zero-shot tracking accuracy on an unseen (species, place) can be *predicted in advance* from properties measurable without any label of the target, and which factors govern that transfer — separately for *finding* the animal and *following* it.

1.3 Hypotheses

H0 (null)

the label-free distances carry no predictive information about pDetA/pAssA — transfer is not distance-driven.

H1 (detection)

pDetA falls with species novelty (taxonomic + visual distance).

H2 (association)

pAssA falls with environment difficulty (scene, day/night/IR, clutter, camera motion).

Failing to reject H0 is itself an informative result — and, as Chapter 4 shows, it is what we find: neither a novelty distance nor a follow-up *intrinsic-difficulty* pivot predicts transfer beyond a trivial animal-size effect. Representational probing (does SAM3 encode taxonomy?) remains the open branch for future work.

1.4 Contributions

To our knowledge this is the first study to *predict*, before the model is run and from label-free distances, the zero-shot transfer of a promptable video tracker — and the first to split that prediction into **detection** (pDetA) and **association** (pAssA). This distinguishes it from label-free performance prediction for classifiers, which estimates a single scalar from the model’s own outputs on the target [1, 6, 10], and from the very recent static-image detector/segmenter estimators [14, 16]; and from camera-trap shift studies, which measure but do not forecast per-site reliability [2, 3, 7]. Concretely we deliver: (i) a rigorous, out-of-sample *negative* result —

under a deliberately powerful test (whole-species and whole-location hold-out with group-cluster-bootstrap significance), label-free before-running distances do *not* predict SAM3’s transfer; (ii) the detection-versus-association decomposition, tested on two near-orthogonal splits — a constructed species hold-out (novelty) and the official location hold-out (environment); and (iii) a *hardened* account of *why* it fails: a follow-up intrinsic-difficulty pivot collapses to the same animal-size confound, so the only label-free property that predicts detection is object size — the very nuisance the analysis controls for — and that positive control confirms the test detects signal when it exists.

1.5 Dissertation outline

Chapter 2 reviews the related work; Chapter 3 describes the data, measurement, distance features, and modelling; Chapter 4 reports the results and out-of-sample validation; Chapter 5 discusses them; Chapter 6 concludes.

Chapter 2

Background and Related Work

This chapter places the dissertation among four bodies of work: promptable segmentation and tracking (the model we study), the tracking metric and its detection/association decomposition (how we score it), label-free prediction of out-of-distribution performance (the statistical idea we extend), and distribution shift in camera-trap vision (the application and the closest empirical neighbours). It closes by stating precisely the gap this work fills.

2.1 Promptable segmentation and tracking

The *Segment Anything* line moved segmentation from a fixed label set to an interactive, promptable formulation, and its successors extended prompting from points and boxes to *concepts*. SAM3 [4] is a promptable model that, given a short text prompt such as "impala", detects, segments, and tracks every matching instance across a video with no training on that category — a *zero-shot*, open-vocabulary detect–segment–track capability. It is exactly this capability that conservation teams now deploy on new species and regions, and it is the (frozen) model whose transfer we predict. SAM3 was released alongside SA-FARI [15], a large multi-animal camera-trap tracking dataset that we use as the substrate for our natural experiment (Section 2.6). We treat SAM3 strictly as a fixed black box: we never fine-tune it, and the only quantity we fit is a small regression *about* its behaviour.

2.2 The tracking metric and its decomposition

Higher Order Tracking Accuracy (HOTA) [9] evaluates multi-object tracking by combining, over a range of localisation thresholds, a *detection* term (**DetA**: were the right objects found?) and an *association* term (**AssA**: were identities kept consistent over time?). SA-FARI scores promptable tracking with a promptable variant, **pHOTA**, that decomposes identically into **pDetA** and **pAssA**. This decomposition is the intellectual core of the dissertation: rather than predict a single scalar, we ask whether *finding* and *following* an animal transfer for *different reasons* — the former with how novel the animal is, the latter with how difficult the scene is — and we therefore model the two components separately throughout.

2.3 Predictable out-of-distribution generalisation

A now-mature literature shows that a model’s performance under distribution shift can be anticipated *without* target labels. Miller et al. [10] observe that out-of-distribution (OOD) accuracy is strongly linearly correlated with in-distribution accuracy across many models and shifts, and Baek et al. [1] show that the *agreement* between a pair of models tracks OOD accuracy closely enough to estimate it from unlabelled data alone. A parallel strand estimates accuracy from a

model’s own statistics on the unlabelled target: average thresholded confidence [6] calibrates a confidence threshold on source data and reads off target accuracy as the mass above it, while Deng and Zheng [5] regress classifier accuracy directly on a *distribution-shift feature* — the Fréchet distance between train and test feature sets — which is the closest methodological cousin to our own “regress the score on a distance” approach.

Two properties characterise essentially all of this work, and both are limitations we depart from. First, it targets **image classifiers** with a single scalar accuracy. The idea has only recently reached structured outputs: Yoo et al. [16] estimate object-detection performance without ground truth from the spatial consistency and confidence reliability of pre-NMS boxes, and reverse classification accuracy [14] predicts per-image segmentation quality by training a reverse model against a labelled reference. These are the closest structured-output predecessors, but both operate on *static single images*, neither concerns identity over time, and none decomposes the estimate into detection versus association. Second, every one of these methods still *runs the model on the target* and inspects its predictions, confidences, or agreement — it estimates the score of an *already-executed* model. We instead forecast from properties knowable *before the model is ever run* on the target.

2.4 Distribution shift in camera-trap and wildlife vision

That camera-trap models degrade at new locations is well established. *Terra Incognita* [2] showed that recognisers excellent at trained camera sites fail sharply at unseen ones; WILDS [8] consolidated this (via iWildCam) as a canonical real-world shift; and PanAf-FGBG [3] demonstrated that behaviour-correlated *backgrounds* causally drive the drop, proposing a latent-space mitigation. These works are the empirical backdrop for our environment axis — but they share two features that mark the gap. They **measure or mitigate** degradation after the fact rather than *predicting* per-site or per-species reliability in advance, and they study **classifiers or behaviour recognisers**, not promptable trackers, scored by a single accuracy with no detection/association split. Most tellingly, a recent unified study of camera-trap recognition over time [7] explicitly names “predicting when zero-shot models will succeed at a new site” as an open question — and stops at naming it. This dissertation builds the predictor that question asks for.

2.5 What should predict transfer

Our candidate predictors are motivated by representation-learning results. Vision foundation models encode taxonomic structure across the tree of life [13], which motivates a *taxonomic* distance and, because appearance and phylogeny are correlated, a *visual* distance; we compute the latter from mask-cropped DINOv2 [11] embeddings, with CLIP [12] as a robustness variant. The camera-trap location results above [2, 3] motivate an *environment* distance from the scene background, and per-video timestamps motivate a *temporal* gap. Crucially, all four are computable without any label of the target animal, and all are measured against a frozen reference set, so no target information can leak into a predictor.

2.6 The SA-FARI benchmark

SA-FARI [15] is a large open multi-animal camera-trap tracking dataset (11,609 densely annotated videos spanning 99 species with a seven-level taxonomy, 741 locations on four continents, and per-video timestamps), released with SAM3 and scored by the official VEval evaluator. It ships with abundant *hard negatives* — prompts for absent species that must return nothing — which we keep, as they drive the false-positive side of detection. Two properties of the shipped split shape our design and merit a precise statement. Inspection of the released annotations

shows the train/test split is disjoint by **location** (test camera sites are unseen) but that its *species are largely shared* across the split; a genuine “unseen species” axis is therefore not directly available from the official split. Because SAM 3 is frozen and zero-shot on *all* of SA-FARI, however, the split reflects nothing the model learned — it is only our choice of a *reference distribution* against which distances are measured. This licenses a species-disjoint *reference/probe partition* of the same videos for the primary (species-novelty) experiment, complemented by the official location split for the environment experiment; the two are near-orthogonal probes of the same data, and their construction is detailed in Chapter 3.

2.7 Summary and gap

Predictable OOD generalisation is established for classifiers scored by a scalar [1, 5, 6, 10] and, very recently, for static detectors and segmenters [14, 16]; camera-trap vision has thoroughly documented, but not forecast, location-driven degradation [2, 3, 7, 8]. No prior work predicts, *before running the model* and from label-free distances, the zero-shot transfer of a **promptable video tracker**, and none separates that prediction into **detection** and **association**. Asking that question rigorously — an out-of-sample test, over SA-FARI’s species-by-place-by-time structure, of whether label-free distances predict SAM 3’s transfer, decomposed into **pDetA** and **pAssA** — is the contribution of this dissertation. As Chapter 4 reports, the honest answer is negative: transfer is not distance-driven.

Chapter 3

Data and Methodology

3.1 Research design and evaluation criteria

This work is a predictive-modelling study: we ask whether label-free *distances* measurable in advance (X) predict a frozen tracker’s zero-shot performance ($Y = \text{pDetA/pAssA}$), and which factors govern that transfer for *finding* versus *following* an animal. Because the target is a bounded score and the observations are grouped by species and location, the analysis is a weighted regression with group-aware, out-of-sample validation rather than a tuned tracker.

Definition of success and the decision rule. The intended deliverable is a *validated, out-of-sample* predictor of pDetA/pAssA from label-free distances, with honest confidence intervals. A distance is taken to predict transfer only if it satisfies **both** criteria: (i) its coefficient’s group-cluster-bootstrap 95% confidence interval excludes zero, *and* (ii) the grouped cross-validated model beats a mean-predictor baseline by a margin that is significant under a paired group-bootstrap. **H1** (detection $\leftarrow$ species novelty) and **H2** (association $\leftarrow$ environment difficulty) are supported only when their respective distances meet this bar on the split that isolates them; if neither does under these deliberately powerful tests, we conclude **H0** — distance explains little variance — which is itself a valid result and motivates the representational-probing direction of Chapter 5.

3.2 Dataset, splits, and the analysis unit

We use SA-FARI [15] throughout. In the released `_ext` annotations, a probe is a (video, text-prompt) pair; the animal’s species and its seven-level taxonomy are resolved through the prompt’s `category_id`, while per-video `location_id` and creation timestamp give place and time. We aggregate scores to a *cell* — a (`category_id`, species, location, time) group keyed by the stable species identity `category_id` — and attach support counts (n_{frames} , n_{masklets} , n_{videos}) used for weighting and as covariates. Hard negatives (prompts for absent species) are kept, as they drive the false-positive side of detection; the one exception is the stratified per-species cap used to make the large pooled pass feasible, under which hard negatives are set aside from inference and analysed instead by the separate false-positive pass (Section 3.7).

Two complementary splits. Inspection of the annotations shows the shipped split is disjoint by *location* but shares species across train and test (Section 2.6); on that split species-distance is ≈ 0 for every cell, so it cannot test species novelty. Because SAM3 is frozen and zero-shot on all of SA-FARI, the split only defines our *reference distribution*, so we are free to repartition the same videos. We therefore use two near-orthogonal experiments and run inference only once (scoring is per-video and split-agnostic):

Split A — species hold-out (primary).

Pooled species are partitioned by a *leave-one-species-out* view, so a held-out species’ distance is to the nearest *other* species; species-distance varies while the environment stays largely familiar, isolating H1. Validated by leave-species-out cross-validation.

Split B — location hold-out (secondary).

The official train/test split, whose unseen locations isolate H2 and whose benchmark comparability carries the Phase-1 score-sanity check.

Split A is a *constructed* hold-out — legitimate because the model is frozen, and reported as such. At the scale we work with, the location split (Split B) yields 1,447 scored cells over 113 taxonomy-bearing category identities (99 species) and 92 unseen test locations: 346 are *positive* cells — the queried species is present, so **pDetA**/**pAssA** are defined — and 1,101 are hard-negative cells, drawn from 1,486 hard-negative probes. The species hold-out (Split A) pools the present species of both official splits into 1,025 scored cells across the 99 species.

3.3 Measuring transfer




Figure 3.1: The pipeline. *Top:* frozen SAM3 is run over the videos and scored by the official evaluator into per-cell detection/association scores (the target Y). *Bottom:* four label-free distances are computed from the videos and ground-truth masks against a frozen reference set (the predictors X). A support-weighted logit GLM relates X to Y and is validated out of sample by holding out whole species and whole locations; SAM3 is never trained.

SAM3 [4] is run frozen over the probe videos (species-specific prompts as the primary condition, a generic “animal” prompt for robustness) and scored with the official **VEval** evaluator — the metric is never re-implemented — giving per-cell **pDetA**/**pAssA**/**pHOTA** [9] with their support. A single inference pass over the union of Split A and Split B probes serves both experiments; predictions are stored per (video, **category_id**) so that the several species-prompts issued on one video do not collide, and the harness is resumable at per-video granularity and never updates a weight. We aggregate two ways and keep them distinct: the *micro* dataset-level **VEval** score (the published SA-FARI operating point, ≈ 0.65 mask-pHOTA) and the per-cell support-weighted *macro* mean over positive cells that is our modelling target. Because the macro mean is a positives-only per-cell average it sits lower (≈ 0.53), and every Gate-1 number we quote is the macro value, labelled as such.

No retraining; validity of the frozen protocol. It bears stating plainly, as it is easily misread: **SAM3 is never retrained or fine-tuned in this work.** The single object we fit is the small support-weighted logistic regression that maps the label-free distances X onto the observed scores Y ; the tracker itself is the public checkpoint, run once, its weights untouched. Reusing that one inference across both splits is therefore sound rather than a leak: because the model is frozen, a cell’s score does not depend on which reference partition we later overlay to compute its distances. This also fixes the benchmark’s status here. SA-FARI is a held-out, *zero-shot* benchmark for the released SAM3: the SAM3 report [4] does not list SA-FARI

among its training sources, and the SA-FARI paper [15] reports the released checkpoint zero-shot (the ≈ 0.65 mask-pHOTA ballpark we reproduce at Gate 1) *alongside* separately fine-tuned rows whose gain we deliberately forgo. We can therefore verify non-membership at the *dataset* level; we cannot verify it frame-by-frame, since SAM 3’s full pretraining corpus is undisclosed. We do not claim otherwise: every distance is measured against the SA-FARI *train reference*, not against SAM 3’s true (unknown) pretraining distribution, and the familiarity proxy (Section 3.5, planned) is the hedge for exactly this gap.

3.4 The frozen reference (leakage firewall)

For each split the reference (“seen”) species, locations, and taxonomy are frozen once; every distance is then computed against that reference only, so no probe information can enter a feature — verified by a unit test that swaps the reference and checks the distances move as expected. We assert the axis that is genuinely disjoint (location for Split B; the held-out species for Split A) and *report* any residual species overlap honestly rather than asserting it away.

3.5 Label-free distance features

All distances are label-free for the target and computed against the frozen reference.

3.5.1 Taxonomic distance

Lowest-common-ancestor tree steps from the probe’s species to the nearest reference species (shared genus = 1, family = 2, order = 3, ...); the leave-species-out reference keeps this non-trivial for shared species.

3.5.2 Visual distance

Each species is fingerprinted by a *prototype* — the re-normalised mean DINOv2 [11] embedding of its *ground-truth-mask-cropped* animals — and the visual distance is the cosine gap from a probe species’ prototype to that of the nearest *other* reference species (leave-one-species-out on Split A), with a CLIP [12] encoder as a robustness check. Using the dataset’s ground-truth masks means this feature needs no SAM 3.

3.5.3 Environment distance

Embedding distance between the animal-masked-out background scene and the nearest reference location, plus an interpretable day/night–infrared flag (from per-channel colour statistics of the frame) and a clutter score given by the variance of a Laplacian over the masked background — a cheap measure of local edge/texture density, not a colour statistic.

3.5.4 Temporal gap

The minimum absolute year gap between a cell’s capture year (from the per-video creation-datetime metadata) and the nearest reference year. On the location split the timestamp is present for few cells and near-constant among them, so the temporal axis is reported there as *untested* rather than as a tested null.

3.5.5 Animal size (confound control)

Mean log ground-truth mask area per species. This is not a hypothesis variable but a nuisance covariate, included to separate genuine visual novelty from the fact that visually-distinctive



Figure 3.2: The two image-based distances on a real camera-trap frame (a collared peccary). **(a)** the frame with its ground-truth mask (orange); **(b)** the mask-cropped animal that the *visual* distance fingerprints (background zeroed); **(c)** the same frame with the animal mean-inpainted, which the *environment* distance fingerprints. The per-video capture date (a metadata field, not the pixels) drives the *temporal* gap. None of these features runs SAM 3.

species tend to be large and thus easier to segment; its role in exposing the one apparent effect as an artefact is Section 3.6.

3.5.6 SAM 3 familiarity proxy (deferred)

Separability of the target species in SAM 3’s own feature space would hedge against the unknown true pretraining set. It is *not implemented in the present work* and is noted as future work in the Limitations; no result depends on it.

3.6 Statistical modelling and hypothesis testing

3.6.1 Per-target regression

Each bounded target (pDetA, pAssA) is modelled *separately* with a support-weighted logit-link (fractional-logit) regression over the standardised distances plus the scene covariates and a $\log(n_{\text{frames}})$ covariate, so that “rare” is never mistaken for “far”. Detection and association are analysed apart because their decomposition is the intellectual core of the project. In practice association is a *near-degenerate* target on this data — most cells contain a single object, so pAssA tracks pDetA almost exactly — and it is therefore reported as such rather than as an independently powered test (Chapter 5).

3.6.2 Honest, group-aware inference

Naive model confidence intervals are invalid here for two reasons: support-weighting by frame count inflates the effective sample size, and the distances are constant within a species or a location, so the hundreds of cells are really only tens of independent groups (pseudo-replication). We therefore report *group-cluster-bootstrap* intervals — resampling whole species and whole locations, refitting, and taking percentile intervals — as the honest confidence interval for every coefficient.

3.6.3 Out-of-sample validation and the significance test

The fitted models are validated *out of sample* by holding out whole groups (leave-species-out on Split A, leave-location-out on Split B) and predicting the held-out group. The out-of-sample gain over a mean-predictor baseline is tested with a *paired group-bootstrap* that resamples whole held-out groups, yielding a confidence interval and a one-sided *p*-value on the gain — the criterion (ii) of Section 3.1.

3.6.4 Variance attribution and confound control

Variance is attributed to each factor by dominance/commonality analysis with a VIF collinearity check — not by four univariate fits — and raw correlations are reported as a model-free cross-check. Where a distance reaches significance, we refit with the animal-size covariate to test whether the effect is a size confound rather than novelty.

3.6.5 The intrinsic-difficulty pivot and the nested decomposition

Because the sole apparent effect proved to be a size confound, we ask a complementary question: is transfer governed by the intrinsic *difficulty* of the target imagery rather than its *novelty*? We promote the label-free scene properties already measured for the environment axis — low-light/IR (`achromatic_fraction`), clutter, and the night/IR flag — from nuisance covariates to predictors of interest, and test them against the same out-of-sample bar. To separate a genuine difficulty effect from the size confound, we fit a *nested* sequence of models — support only; size only; difficulty only (no size); difficulty plus size — and compare their out-of-sample gains, so that any apparent difficulty effect that is really size is exposed as such (Chapter 4, Table 4.4). We also report each scene property’s *cell-level* correlation alongside its location-level value, because aggregation inflates correlations (the ecological-correlation fallacy).

Detecting a signal, not asserting its absence. These tests are demonstrably able to detect a real effect: the animal-size covariate not only has an interval excluding zero but *validates out of sample* at the same whole-group hold-out bar the distances fail (Chapter 4), and the estimation and cross-validation machinery is unit-tested to recover planted effects on synthetic data. A null under a test that fires for size is therefore evidence of *absence of signal*, not of *absence of power*.

3.6.6 Design-robustness checks

To show the null is not an artefact of a single design choice, we re-run the entire hold-out procedure with one component of the pipeline swapped at a time: the visual encoder (DI-NOv2 → CLIP), the crop (mask-cropped animal → a whole-frame bounding box that keeps the background), the prompt (species-specific → a generic “animal”), and the visual-distance functional itself (nearest-prototype cosine → a distributional Fréchet/MMD distance between the target and reference embedding sets, after AutoEval [5]). Each variant is fitted and validated at the identical whole-group hold-out bar of Section 3.1, with significance judged after a Bonferroni correction over the variant family; the outcomes are reported in Chapter 4 (Table 4.11).

3.7 Detection precision: false positives on hard negatives

The pDetA model conditions on the animal being present, so it measures recall-given-present, not hallucination. On hard negatives (the queried species is absent) a correct run returns no masklet; we read the prediction directly, flag any returned masklet above a score threshold as a false positive, attach each species’ taxonomic and visual distance, and test — at the species level, to avoid pseudo-replication — whether the hallucination rate rises with novelty.

3.8 Independent replication on external datasets

Every SA-FARI result reuses one frozen inference pass on one dataset, so the two splits are near-orthogonal but not statistically independent. To test whether the before-running null is a property of the transfer problem rather than of SA-FARI, we re-run the *entire* pipeline — unchanged frozen SAM 3, the same official scorer, the same distances, and the same leave-species-out cross-validation and paired group bootstrap — on two independent tracking datasets. The

abstraction boundary that makes this cheap is the annotation schema, not the loader: an offline adapter converts each dataset into the SA-Co JSON the loader and evaluator already read, so nothing downstream changes. Only the small logit GLM is ever fitted; SAM3 is frozen throughout.

MammAlps (box-only, small). An Alpine camera-trap tracker: 5 species over 3 sites / 9 cameras, box-only ground truth. The adapter emits each box as a filled-rectangle RLE (so the mask-cropped visual, size and environment features run) together with the native box, and the scorer reads box-HOTA rather than mask-HOTA (against a filled rectangle a tight predicted mask has near-zero IoU, so box-IoU is the honest metric). The 30 fps clips are subsampled to 3 fps and capped per (species, camera) cell; the per-camera site metadata supports both leave-species-out and leave-camera-out. Only three of the four distances are testable (temporal is degenerate over the six-week window; taxonomic is near-constant across five species).

BURST (mask-native, many-species). The animal subset of the BURST validation split (built on TAO): 41 animal species across 132 video $\times$ class cells, with native COCO-RLE masks scored by mask-HOTA. The animal categories are selected from BURST’s 482 LVIS classes by a WordNet filter (hyponyms of *animal.n.01*) and mapped to a seven-level biological taxonomy where resolvable (25/41 species full, the rest partial $\rightarrow$ NaN). BURST carries no site or timestamp metadata, so the cell key is the per-video sequence and leave-species-out is the only valid grouping; the association target is genuinely non-degenerate here, so both components are tested. Masks stream from the RWTH server and TAO frames from the Hugging Face mirror by HTTP range, so only the annotated frames of the capped animal videos are fetched.

Interpreting the replications. Both datasets are held to the exact same out-of-sample bar as SA-FARI, and both include the animal-size confound control and the after-running confidence features as positive controls — a fired control shows the test is alive at that sample size, and its failure to fire (as on the 20-cell MammAlps) marks the design as underpowered rather than the hypothesis as confirmed. We report convergence across datasets, not independent proof: BURST detects only a large between-species predictor, so it cannot exclude a small before-running distance effect (§4.9).

3.9 Reproducibility, implementation, and constraints

The model is frozen and scored by the official evaluator; the seen set is frozen before any distance is computed; hard negatives are retained; and all inference is single-GPU and inference-only. The pipeline is config-driven (every constraint is a toggle, so an ablation flips config, not code) and covered by hermetic tests; every reported number is regenerated from the repository, and split definitions, seeds, and reference sets are persisted so both experiments are reproducible end to end. The engineering safeguards that keep the measurement honest — memory-lean feature assembly, per-(video, `category_id`) prediction keying, trimming the ground truth to the scored videos before evaluation, and the size-confound control — are catalogued in the appendix, each framed as a validity safeguard.

Chapter 4

Results and Evaluation

4.1 Measurement

Zero-shot SAM 3 on the SA-FARI test split reproduces the published ballpark (Gate 1); the support-weighted per-cell aggregate is in Table 4.1.

Table 4.1: Zero-shot SAM 3 on the SA-FARI test split — support-weighted mean over 346 scored cells (of 1447).

Metric (mask)	Value
pDetA	0.527
pAssA	0.546
pHOTA	0.533

4.2 The fitted law

Table 4.2 gives the standardised logit-GLM coefficients for detection and association.

Table 4.2: Standardised logit-GLM coefficients with *group-cluster-bootstrap* 95% CIs (resampling whole species / locations, so the interval respects pseudo-replication and the support weighting), for detection (pDetA) and association (pAssA). Every distance’s CI spans zero.

Factor	pDetA coef. [CI]	pAssA coef. [CI]
Taxonomic distance	-0.362 [-3.87, 0.09]	-0.368 [-3.40, 0.08]
Visual distance	-0.144 [-0.50, 0.14]	-0.122 [-0.49, 0.18]
Environment distance	0.042 [-0.19, 0.24]	0.013 [-0.23, 0.21]
Clutter	-0.066 [-0.38, 0.29]	-0.087 [-0.41, 0.28]
Night / IR	-0.405 [-0.94, 0.11]	-0.462 [-1.02, 0.11]
log(support)	-0.216 [-0.51, 0.17]	-0.224 [-0.53, 0.18]




Figure 4.1: Standardised GLM coefficients with group-cluster-bootstrap 95% CIs, for detection (pDetA) and association (pAssA). *Every distance’s interval crosses zero (grey) — none is distinguishable from no effect.*

4.3 Variance partitioning and the intrinsic-difficulty decomposition

A dominance (LMG/Shapley) analysis confirms the distances carry almost nothing: together the four explain only $\approx 3\%$ of the variance in detection, no single factor dominates, and every VIF is ≈ 1 (Table 4.3) — so the null is not an artefact of collinearity masking a real effect. Because the

Table 4.3: Variance partition for detection (pDetA): each distance’s unique share of R^2 (LMG/Shapley) and its VIF. The four distances together explain only $\approx 3\%$ of the variance, no factor dominates, and every VIF ≈ 1 — the null is not masked by collinearity.

Factor	unique R^2	share	VIF
Taxonomic	0.014	48%	1.06
Visual	0.012	39%	1.06
Environment	0.004	13%	1.00
Temporal	0.000	0%	–

one apparent effect (visual distance, below) turned out to be an animal-size confound, we pivot the hypothesis from *novelty* to intrinsic *difficulty*: we promote the label-free scene properties — low-light/IR (`achromatic_fraction`), clutter, and the night/IR flag, each computed per location from the frames with no SAM3 run and no target label — from nuisance covariates to predictors of interest, and test them at the same bar. Table 4.4 decomposes the outcome. Taken *with* an animal-size covariate each difficulty signal *appears* to validate (leave-species-out and leave-location-out $\Delta\text{MAE} \approx +0.017$, both significant), but a nested decomposition shows the gain is entirely the size covariate: log-area *alone* carries it ($\Delta\text{MAE} = +0.017$, $p = 0.001/0.014$), while low-light *alone* does *not* validate ($p = 0.12/0.23$) and adds only $\approx +0.002$ over size. The low-light effect is also far weaker than it first looks: the location-level night/IR–pDetA correlation of $r = -0.377$ is an aggregation artefact (the ecological-correlation fallacy) — at the cell level it is only -0.13 , and clutter’s cell-level correlation is 0.00 . The intrinsic-difficulty pivot is therefore **also null**: the only label-free property that predicts detection is animal size.

Table 4.4: Nested size decomposition of the intrinsic-difficulty test (detection, 346 cells). Each entry is the out-of-sample MAE gain over a mean-predictor baseline, Δ [95% CI] (p), under leave-species-out and leave-location-out. The entire gain sits in the size covariate (log-area); low-light *alone* does not validate — the difficulty signals only pass the bar by riding on size.

Model	Leave-species-out Δ [CI] (p)	Leave-location-out Δ [CI] (p)
log(support) only	+0.001 [-0.007, +0.008] (0.432)	+0.002 [-0.005, +0.008] (0.291)
log(area) only [<i>size</i>]	+0.015 [+0.002, +0.028] (0.014)	+0.017 [+0.009, +0.025] (0.001)
Low-light only (no size)	+0.003 [-0.006, +0.011] (0.233)	+0.004 [-0.003, +0.011] (0.123)
Low-light + size	+0.017 [+0.004, +0.032] (0.011)	+0.020 [+0.011, +0.030] (0.000)

4.4 Out-of-sample validation

Table 4.5 reports whole-group hold-out error against a mean-predictor baseline. This is the headline result and it is a **null**: on both the location and the species schemes, and for both targets, no distance model beats simply predicting the mean — every Δ MAE is within noise of zero and no gain is significant under the paired group-bootstrap. **This is not an underpowered test.** The same GLM and grouped cross-validation, on the same cells, *do* detect a real signal when one exists: animal size (log-area) validates out of sample at the identical bar (Δ MAE = +0.017, $p < 0.017$ on both schemes; Table 4.4). A test that fires for size yet stays flat for the distances is evidence of a genuine *absence* of before-running label-free signal, not of low power.

Table 4.5: Out-of-sample error: whole-group hold-out (leave-species-out / leave-location-out) vs a mean predictor. Δ is the MAE gain over baseline; its bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. No hold-out’s gain is significant.

Hold-out	Target	n	MAE	Baseline	Δ [95% CI]	p
location	pAssA	346	0.289	0.295	+0.007 [-0.005, +0.018]	0.13
location	pDetA	346	0.293	0.298	+0.004 [-0.007, +0.015]	0.23
species	pAssA	346	0.289	0.295	+0.006 [-0.006, +0.017]	0.16
species	pDetA	346	0.293	0.298	+0.004 [-0.006, +0.015]	0.24




Figure 4.2: Out-of-sample predictions vs the truth for detection (pDetA), on held-out species and locations. The predictions collapse onto the mean baseline (orange) instead of tracking the diagonal (dashed) — the distances carry no out-of-sample predictive power.

4.5 Species hold-out (H1): novelty and the size confound

The primary novelty test — leave-one-species-out over the pooled train+test species (99 species, 1025 scored cells) — is in Tables 4.6–4.8. Out-of-sample prediction is not significant on any hold-out (Table 4.8). The only coefficient whose CI excludes zero, *visual distance* (Table 4.7), is positive — meaning more-novel species are detected *better*, the reverse of what H1 predicts — and is a **confound**: visually-distinctive species are larger (visual $\leftrightarrow$ log-area $r = 0.22$) and larger animals score higher (log-area $\leftrightarrow$ pDetA $r = 0.30$, the strongest raw correlate of the score). Adding an animal-size covariate shrinks the visual coefficient from +0.37 to +0.22 (CI now spanning zero) while size itself is the significant driver, so “visual novelty helps” was really “bigger animals are easier”. Taxonomic novelty is flat ($r \approx 0.00$). Novelty does not predict transfer.

Table 4.6: Zero-shot SAM3 on the pooled train+test set — support-weighted mean over 1025 scored cells (of 2126).

Metric (mask)	Value
pDetA	0.536
pAssA	0.597
pHOTA	0.559

Table 4.7: Standardised logit-GLM coefficients with *group-cluster-bootstrap* 95% CIs (resampling whole species / locations, so the interval respects pseudo-replication and the support weighting), for detection (pDetA) and association (pAssA).

Factor	pDetA coef. [CI]	pAssA coef. [CI]
Taxonomic distance	0.202 [-0.09, 0.43]	0.209 [-0.06, 0.42]
Visual distance	0.369 [0.11, 0.67]	0.271 [0.04, 0.54]
Environment distance	0.043 [-0.15, 0.17]	0.023 [-0.15, 0.13]
Clutter	-0.284 [-0.52, 0.02]	-0.230 [-0.46, 0.06]
Night / IR	-0.286 [-0.77, 0.22]	-0.439 [-0.92, 0.06]
log(support)	-0.020 [-0.14, 0.16]	0.051 [-0.06, 0.23]

Table 4.8: Out-of-sample error: whole-group hold-out (leave-species-out / leave-location-out) vs a mean predictor. Δ is the MAE gain over baseline; its bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. No hold-out’s gain is significant.

Hold-out	Target	n	MAE	Baseline	Δ [95% CI]	p
location	pAssA	1025	0.270	0.270	-0.000 [-0.006, +0.005]	0.52
location	pDetA	1025	0.274	0.274	-0.001 [-0.007, +0.006]	0.58
species	pAssA	1025	0.273	0.270	-0.003 [-0.014, +0.008]	0.74
species	pDetA	1025	0.277	0.274	-0.004 [-0.015, +0.009]	0.74




Figure 4.3: The one label-free property that predicts detection: animal size. Per-cell pDetA rises with the log ground-truth mask area ($r = +0.31$) — the confound behind the apparent visual-distance effect.

4.6 False positives / hallucination

The detection score conditions on the animal being present, so it measures recall, not hallucination. The 1486 *hard negatives* (queries for absent species) ask the complementary question, and are the closest the study comes to a validated before-running signal. SAM3 returns a masklet anyway about **9.8%** of the time. This hallucination rate does *not* rise with taxonomic novelty ($r \approx +0.06$, n.s.); it *falls* with visual distinctiveness ($r \approx -0.33$), and this survives an animal-size control (partial -0.37) — a genuine, size-independent, label-free correlate, though its direction (“distinctive = safer”) is the *opposite* of “novelty hurts”. Crucially it is a **correlate, not a predictor**: predicting a held-out species’ hallucination rate from visual distance beats a mean baseline by only $+0.008$ MAE and *just misses* out-of-sample significance ($p = 0.053$; Table 4.9). Even the study’s best signal does not clear the bar.

Table 4.9: False positives on the 1486 hard negatives (113 species): the overall hallucination rate is 9.8%. More visually distinctive species hallucinate *less* — a real, size-independent correlate — but it does not validate out of sample (leave-species-out $p = 0.053$), so it is a correlate, not a predictor.

Test (species-level)	statistic	95% CI	verdict
FP rate vs taxonomic distance	+0.06	[−0.17, +0.23]	n.s.
FP rate vs visual distance	−0.33	[−0.43, −0.23]	significant
visual, size-controlled (partial)	−0.37	[−0.47, −0.28]	significant
FP rate vs animal size	+0.13	[−0.06, +0.31]	n.s.
leave-species-out OOS gain	+0.008	[−0.002, +0.017]	$p = 0.053$

4.7 Ablations and robustness

Factor ablation. Refitting detection with each distance dropped in turn barely moves the fit or the (null) out-of-sample gain (Table 4.10): no single distance carries signal. The only design that changes anything is adding the animal-size covariate, which alone lifts the out-of-sample gain to significance ($\Delta\text{MAE} \approx +0.02$, $p \leq 0.02$) — again consistent with size being the sole real correlate. Removing the support covariate does not rescue a distance effect either. The variance partition (Table 4.3) tells the same story from the other direction.

Table 4.10: Factor ablation for detection (pDetA): McFadden pseudo- R^2 and the out-of-sample MAE gain over a mean baseline under leave-species-out (LSO) and leave-location-out (LLO). Dropping any single distance barely changes the fit or the (null) out-of-sample gain — no factor carries signal; only adding the animal-size covariate moves the score.

Design	pseudo- R^2	LSO Δ (p)	LLO Δ (p)
full (4 distances)	0.055	+0.004 ($p=0.24$)	+0.004 ($p=0.23$)
drop taxonomic	0.039	+0.000 ($p=0.47$)	+0.000 ($p=0.49$)
drop temporal	0.055	+0.004 ($p=0.24$)	+0.004 ($p=0.23$)
drop visual	0.048	+0.004 ($p=0.22$)	+0.004 ($p=0.28$)
drop environment	0.055	+0.006 ($p=0.15$)	+0.005 ($p=0.17$)
+ animal size	0.119	+0.018 ($p=0.02$)	+0.020 ($p=0.00$)
− support covariate	0.035	−0.002 ($p=0.72$)	−0.003 ($p=0.73$)

Design robustness. The null is not an artefact of the encoder, the crop, the prompt, or the distance functional (Table 4.11). Re-running the full leave-species-out and leave-location-out

procedure with a CLIP rather than DINOv2 encoder, a whole-frame rather than mask-cropped embedding, and a distributional Fréchet or MMD rather than nearest-prototype distance each leaves the out-of-sample gain indistinguishable from the baseline (all $\Delta\text{MAE} \leq 0.007$, $p \geq 0.10$, every 95% CI spanning zero). A generic “animal” prompt is included for completeness: prompting non-specifically collapses per-species detection to a near-constant, mostly-zero target (mean $\text{pDetA} \approx 0.14$, 73% of cells exactly zero),¹ so the distance model cannot beat — and, fitting a degenerate target, underperforms — a mean predictor. Across every variant the label-free distances fail identically, confirming the null is a property of the transfer problem, not of one modelling choice.

Table 4.11: Design-robustness sweep (detection). Each row re-runs the primary four-distance model with one design choice swapped; entries are the out-of-sample MAE gain over a mean-predictor baseline, Δ [95% CI] (p), under leave-species-out and leave-location-out, with significance judged after Bonferroni ($m = 5$). The null holds throughout: no encoder, crop, prompt, or distributional distance makes the label-free distances beat the mean.

Variant	Leave-species-out Δ [CI] (p)	Leave-location-out Δ [CI] (p)
baseline (DINOv2, mask-crop, species prompt) (<i>baseline</i>)	+0.004 [-0.006, +0.015] (0.237)	+0.004 [-0.007, +0.015] (0.237)
CLIP encoder	+0.004 [-0.006, +0.014] (0.247)	+0.004 [-0.007, +0.014] (0.247)
whole-frame (bbox + background)	+0.002 [-0.008, +0.013] (0.349)	+0.002 [-0.008, +0.013] (0.349)
generic “animal” prompt	-0.346 [-0.381, -0.309] (1.000)	-0.343 [-0.363, -0.323] (1.000)
Frechet distributional distance	+0.007 [-0.005, +0.019] (0.120)	+0.007 [-0.004, +0.018] (0.120)
MMD distributional distance	+0.002 [-0.008, +0.012] (0.374)	+0.002 [-0.009, +0.013] (0.374)

4.8 Independent replication (MammAlps)

Every result so far reuses one frozen inference pass on one dataset, so the species and location experiments are near-orthogonal overlays rather than statistically independent replications. To test the null on genuinely independent data we re-ran the whole pipeline on **MammAlps** — an Alpine camera-trap tracker with a different continent, species pool, and (uniquely) real per-camera environment metadata. MammAlps ships box-only ground truth, so an offline adapter emits each box as a filled-rectangle mask and the scorer reads bbox-HOTA (`eval.prefer_bbox`); mask-IoU against filled rectangles is near-zero (dataset mask mAP ≈ 0.002), so box-IoU is the honest metric. The run yields 20 cells over 5 species (red deer, roe deer, hare, wolf, and a single-cell fox) and 9 cameras. Zero-shot SAM3 is non-degenerate here ($\text{pHOTA} \approx 0.44$, Table 4.12); as on SA-FARI, association tracks detection almost exactly ($\text{corr}(\text{pDetA}, \text{pAssA}) = 0.999$, 13/20 cells identical), so the replication is a detection story.

The distance null holds, but as a consistency check — not confirmation. The pre-registered label-free distance model (taxonomic + visual + environment) does not predict detection transfer out-of-sample under either hold-out (leave-camera-out $\Delta\text{MAE} = +0.067$, 95% CI [-0.043, +0.144], $p = 0.099$; leave-species-out +0.014, $p = 0.410$; both intervals span zero, Table 4.13), directionally consistent with the SA-FARI null. We read this as a consistency check rather than independent confirmation, because at $n = 20$ the out-of-sample test is demonstrably underpowered: a *genuinely* predictive per-cell feature — SAM3’s own mean confidence, in-sample $r = +0.57$ ($p = 0.009$) — itself fails to validate ($\Delta\text{MAE} = -0.03 / -0.05$, $p > 0.9$). A real signal that cannot clear the bar shows the instrument, not the hypothesis, is the limiting factor here.

¹A minority of clips were additionally skipped for GPU memory on the 16 GB card used for this pass; they are scored as misses and add noise but do not drive the collapse, which is intrinsic to the non-specific prompt.

Table 4.12: Zero-shot SAM3 on MammAlps — support-weighted mean over the 20 probe cells (MammAlps ships no separate reference split; all clips are probe). Scored with **bbox**-HOTA (`eval.prefer_bbox`): the ground truth is box-only, so mask-IoU is near-zero (dataset mask mAP ≈ 0.002) and box-IoU is the honest metric. The score is non-degenerate — SAM3 genuinely finds Alpine mammals zero-shot.

Metric (bbox-HOTA)	Value
pDetA	0.440
pAssA	0.443
pHOTA	0.441

Table 4.13: MammAlps independent replication — out-of-sample error of the pre-registered label-free distance model (taxonomic + visual + environment) under whole-group hold-out, against a mean predictor. Hold-out *camera* is leave-one-camera-out (the environment scheme); *species* is leave-one-species-out. Δ is the MAE gain over baseline; the bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out groups. Every interval spans zero on both targets and both schemes: the label-free distances do not predict transfer, directionally consistent with the SA-FARI null. $n = 20$ cells (5 species, 9 cameras); a single-cell fox fold and near-constant taxonomic distance leave the species scheme underpowered.

Hold-out	Target	n	MAE	Baseline	Δ [95% CI]	p
camera	pDetA	20	0.210	0.277	+0.067 [−0.043, +0.144]	0.099
camera	pAssA	20	0.205	0.274	+0.069 [−0.040, +0.145]	0.095
species	pDetA	20	0.263	0.277	+0.014 [−0.218, +0.162]	0.410
species	pAssA	20	0.255	0.274	+0.018 [−0.215, +0.165]	0.410

The one feature that “validates” is a confound, not a finding. Visual distance alone beats the mean baseline on both schemes ($\Delta\text{MAE} = +0.078/ +0.080$, $p \approx 0.01$), but this is spurious and is reported only to be dismissed. Its association with detection is *positive* ($r = +0.636$, $p = 0.003$) — the opposite direction to the novelty→failure hypothesis (more-novel animals are detected *better*, not worse), driven by the two large, high-footage deer species SAM3 detects best; it is confounded with support ($r = +0.60$ with n_{frames}); and because visual distance is species-constant with tied values (red = roe = 0.40, hare = wolf = 0.15), its leave-species-out “validity” reduces to a twin species standing in for the held-out one — exactly the species-constant-feature artefact the double cross-validation is designed to reject. Removing the two rare, low-support species that anchor it (hare, wolf; 4/20 cells) reverses the correlation to $r = -0.33$ (n.s.) and erases the gain on both schemes. This is the same visual/size confound diagnosed on SA-FARI, and it does not support H1.

Scope. MammAlps exercises only three of the four distances (temporal and the familiarity proxy are degenerate over the six-week window; taxonomic distance is near-constant across five species), scores a different target regime (box-HOTA at 3fps), and is underpowered by construction (a single-cell fox fold). It therefore reinforces — and does not contradict — the SA-FARI null, but cannot serve as independent confirmation of any positive result. A larger, mask-native second dataset remains the route to that.

4.9 Independent replication (BURST)

The larger, mask-native dataset MammAlps pointed to is BURST. Its animal subset — the animal LVIS classes of the validation split, streamed mask-native and prompted by class name — yields 132 video $\times$ class cells across 41 species, an order of magnitude more leave-species-out

groups than MammAlps, in a distinct regime (internet/driving/movie video, not camera traps). Zero-shot SAM3 is again non-degenerate (pHOTA 0.635, Table 4.14), and here the association target is genuine ($\text{corr}(\text{pDetA}, \text{pAssA}) = 0.76$; of the 71 multi-object cells, none has pAssA exactly equal to pDetA), so both components are testable.

A third convergent null. Under leakage-free leave-species-out cross-validation the before-running distances again fail to predict detection transfer: the combined distance model improves out-of-sample MAE by only +0.002 ($p = 0.374$), with visual (+0.003, $p = 0.257$), taxonomic (full-taxonomy subset -0.006 , $p = 0.952$) and environment (-0.001 , $p = 0.747$) all individually null, and the association specs null likewise (Table 4.15). The result is robust to dropping the 19 single-cell species and to a size control. The lone near-signal, visual distance, is the familiar artefact: its correlation with detection is +0.17 — in the *opposite* direction to the novelty hypothesis (more-novel animals detected better), and a size confound (partial correlation given log-area +0.08, $p = 0.35$). Three datasets — SA-FARI, MammAlps and BURST — now converge on the same before-running distance null.

Table 4.14: Zero-shot SAM3 on the BURST val animal subset — support-weighted mean over the 132 probe cells (41 species). Mask-native ground truth, scored with mask-HOTA (dataset mask mAP 0.49, bbox mAP 0.51). The score is non-degenerate and the association target is genuine ($\text{corr}(\text{pDetA}, \text{pAssA}) = 0.76$, 71/132 multi-object cells).

Metric (mask-HOTA)	Value
pDetA	0.607
pAssA	0.688
pHOTA	0.635

Table 4.15: BURST independent replication — out-of-sample error under leakage-free leave-species-out cross-validation (41 species, 132 mask-native cells), against a mean predictor. Δ is the MAE gain over baseline for pDetA; the bracketed 95% CI and one-sided p come from a paired bootstrap resampling whole held-out species. Every before-running distance is null (CI spans zero); the after-running confidence controls fire, but through a within-species channel a species-constant distance cannot use (collapsing `conf_mean_score` to species means drops it to +0.004, $p = 0.21$), so they establish pipeline liveness, not power for a distance. BURST has no shared locations, so leave-species-out is the only valid scheme.

Predictor	ΔMAE	95% CI	p
distances (taxonomic + visual + environment)	+0.002	[−0.007, +0.011]	0.374
visual only	+0.003	[−0.005, +0.013]	0.257
taxonomic only (full-taxonomy subset, $n=82$)	−0.006	—	0.952
environment only	−0.001	[−0.005, +0.003]	0.747
animal size (<code>log_area</code>) [control]	+0.005	[−0.003, +0.014]	0.143
<code>conf_mean_score</code> [after-running control]	+0.035	[+0.011, +0.063]	0.004
<code>conf_atc</code> [after-running control]	+0.029	[+0.004, +0.055]	0.012

Powered, but not proof. On the same cells the after-running confidence controls do fire (`conf_mean_score` +0.035, $p = 0.004$; `conf_atc` +0.029, $p = 0.012$), a real contrast with MammAlps, where nothing validated: at 132 cells the pipeline is demonstrably alive, not a dead or misconfigured test. This does not, however, license calling the distance null fully powered. The confidence win runs mostly through a *within*-species channel — collapsing `conf_mean_score` to species means, as a species-constant distance effectively is, drops its gain to +0.004 ($p = 0.21$) — whereas a distance can only act *between* species, where the correlations are at most +0.37

(visual, and in the opposite direction to novelty) against confidence’s +0.50. The design thus detects a large between-species predictor but is underpowered for a small one. We therefore read BURST as a third convergent null — no moderate-or-large before-running distance signal, in a better-powered mask-native regime — rather than positive proof, with the only validated predictor across all three datasets remaining SAM 3’s own after-running detection confidence, a mechanistically near-circular, non-novel signal.

4.10 Cross-dataset synthesis

Table 4.16 collects every feature tested, its numerical out-of-sample result, and the dataset it was measured on. The pattern is consistent across all three: no before-running label-free distance — taxonomic, visual, or environment, alone or combined — predicts detection or association transfer out of sample; every gain over a mean predictor is within bootstrap noise of zero. The only near-signal, visual distance, appears in the *opposite* direction to the novelty hypothesis (more-novel animals detected better, not worse) and dissolves under an animal-size control wherever it is examined. The single feature that validates is not a before-running distance at all but SAM 3’s own *after-running* detection confidence, and even that decisively only on the large SA-FARI sample; the same animal-size confound is the only before-running quantity that ever reaches significance. The contribution is this robust, thrice-replicated null, honestly bounded: consistent across a camera-trap and an internet-video dataset, and powered enough on SA-FARI and BURST to exclude a moderate effect, but not a small one.

Table 4.16: What was tested, the numerical out-of-sample result, and on which dataset. Result is the leave-species-out gain in MAE over a mean predictor (ΔMAE , with one-sided bootstrap p) unless noted; positive and significant means the feature predicts transfer. The before-running label-free distances are null on all three datasets; the only feature that validates is SAM3’s own *after-running* confidence (and only decisively on SA-FARI — on BURST it fires through a within-species channel a distance cannot use, and on MammAlps the design is too small for anything to validate). “Opposite direction” means the association runs the reverse of what novelty predicts — more-novel animals are detected *better*, not worse — so it cannot support the hypothesis.

Feature tested	Result (leave-species-out)	Dataset
<i>Before-running label-free distances (the research question)</i>		
Combined distances (tax. + vis. + env.)	$\Delta\text{MAE} -0.004$ ($p=0.74$), n.s.	SA-FARI (1025 cells, 99 sp.)
Combined distances (tax. + vis. + env.)	$\Delta\text{MAE} +0.014$ ($p=0.41$), n.s.	MammAlps (20 cells, 5 sp.)
Combined distances (tax. + vis. + env.)	$\Delta\text{MAE} +0.002$ ($p=0.37$), n.s.	BURST (132 cells, 41 sp.)
Visual distance alone	coef. -0.14 $[-0.50, 0.14]$, n.s.	SA-FARI
Visual distance alone	$\Delta\text{MAE} +0.080$ ($p=0.01$) — <i>opposite direction; size confound</i>	MammAlps
Visual distance alone	$\Delta\text{MAE} +0.003$ ($p=0.26$) — opposite direction	BURST
Taxonomic distance alone	coef. -0.36 $[-3.87, 0.09]$, n.s.	SA-FARI
Taxonomic distance alone	$\Delta\text{MAE} -0.08$ ($p=0.94$), n.s.	MammAlps
Taxonomic distance alone (full-tax. $n=82$)	$\Delta\text{MAE} -0.006$ ($p=0.95$), n.s.	BURST
Environment distance alone	coef. $+0.04$ $[-0.19, 0.24]$, n.s.	SA-FARI
Environment distance alone	$\Delta\text{MAE} -0.001$ ($p=0.75$), n.s.	BURST
Temporal gap	n.s. (present); degenerate elsewhere	SA-FARI (MammAlps/BURST: all-NaN)
<i>Confound control</i>		
Animal size (log-area)	$\Delta\text{MAE} +0.015$ ($p=0.01$) — validates	SA-FARI
Animal size (log-area)	$\Delta\text{MAE} -0.015$ ($p=0.65$), n.s.	MammAlps
Animal size (log-area)	$\Delta\text{MAE} +0.005$ ($p=0.14$), n.s.	BURST
<i>After-running (reads SAM3’s own outputs — not before-running)</i>		
Detection confidence (ATC)	$\Delta\text{MAE} +0.143$ ($p\approx 0$) — validates	SA-FARI
Detection confidence (ATC)	$\Delta\text{MAE} +0.029$ ($p=0.01$) — fires (liveness only)	BURST
Detection confidence (mean score)	$\Delta\text{MAE} -0.05$ ($p=0.96$) — fails (underpowered)	MammAlps
<i>Detection precision (hard negatives)</i>		
Hallucination rate vs. visual distance	$r = -0.33$; LSO $\Delta\text{MAE} +0.008$ ($p=0.053$), marginal	SA-FARI

Chapter 5

Discussion

5.1 Interpretation

Neither hypothesis is supported. H1 (detection falls with species novelty) and H2 (association falls with environment difficulty) both fail the pre-registered bar: no label-free distance beats a mean predictor out of sample, and the one coefficient that reached significance — visual distance, running in the opposite direction to the hypothesis (more-novel species detected better, not worse) — is an animal-size artefact. The follow-up intrinsic-difficulty pivot fares no better: low-light, clutter and night/IR predict detection only by proxying animal size, which alone carries the out-of-sample gain. So the honest reading is H0: transfer is *not* distance-driven, and the only label-free property that predicts detection is object size — the very quantity introduced as a nuisance control.

The detection–association decomposition, the intellectual core of the design, has a second, quieter finding: association is a near-degenerate target on this data. Because most cells contain a single object, pAssA tracks pDetA almost exactly, leaving little identity-specific variance for any feature to predict. The decomposition is therefore honestly a detection story; we report the association null with those divergence figures as its explanation rather than as a separate failed search.

5.2 Relation to prior label-free evaluation

The null sits against a literature of *positive* results, and two axes separate every one of those results from ours — which is why the disagreement is only apparent. The first axis is the **signal**. Accuracy- and Agreement-on-the-Line [1, 10], ATC [6] and AutoEval [5] all read the model’s *own behaviour on the target* — its confidence, its agreement across a model pair, or (for AutoEval) a feature-distribution distance computed after passing the target images through the network. Those are *after-running* estimators. We instead forecast from *external*, label-free distances known *before* a single frame is processed — the strictly harder, and for a practitioner deciding whether to spend compute, the more useful regime. The second axis is the **output**. Every predecessor targets a *scalar* classifier accuracy [1, 5, 6, 10] or the quality of a *static* detector/segmenter [14, 16]; none concerns identity maintained over time, and none is decomposed. Ours is the first study to ask the question of a promptable *video tracker*, split into detection (pDetA) and association (pAssA). AutoEval [5] is the closest methodological cousin — it, too, regresses a performance score on a distributional feature distance — and it is precisely the recipe whose strong classifier-level correlation does *not* reproduce once the distance is measured before running and the target is a tracker. Our result therefore does not contradict these positives; it extends the question to a regime where the label-free-distance shortcut stops working, and reports honestly that it does.

5.3 Why label-free distances do not yield a deployment estimator

The field problem that motivated this work — telling a practitioner, before spending compute, whether SAM3 is trustworthy for a new (species, place) — is not solved by label-free distances. That is a negative but useful answer: the intuition that “far-from-training implies unreliable” does not hold for a promptable video tracker, so a deployment gate cannot be built from taxonomic, visual, environmental or temporal distance to a reference set alone. The one predictive label-free signal, animal size, is too coarse and too obvious to serve as such a gate. What the negative result does establish is a rigorous baseline: any future before-running estimator must clear the same whole-group hold-out bar that these well-motivated distances failed.

5.4 Limitations and threats to validity

Three limitations bound the claim. First, distances are measured against the SA-FARI train split, *not* SAM3’s true pretraining corpus (Meta’s SA-Co), which is undisclosed; the null is scoped to distance-from-SA-FARI, and a coverage-anchored proxy for the real corpus remains untested (the deferred familiarity proxy is one such hedge). Second, the analysis is a single dataset and a single frozen backbone, and — since one inference pass is reused across both splits — the species and location experiments are two overlays on the *same* scored cells, not independent replications; they give near-orthogonal but not statistically independent evidence. We take two steps against this by re-running the whole pipeline on independent datasets. On MammAlps (§4.8) the distance null holds but the 20-cell design is underpowered; on the larger, mask-native BURST (§4.9, 132 cells / 41 species) it holds again and better-powered, giving a third convergent null. Neither is positive proof — BURST detects only a large between-species predictor and remains underpowered for a small species-constant distance effect — so we report convergence, not independent confirmation of any positive. Third, the effective sample is tens of independent species/locations, not hundreds of cells — which is why inference is group-clustered throughout, and why aggregate correlations (e.g. the location-level night/IR effect) are reported alongside their much weaker cell-level values to avoid the ecological-correlation fallacy. The open branch this work does not pursue is representational probing — whether SAM3’s own features encode the species.

5.5 Ethical considerations

The data are anonymised camera-trap footage released under a non-commercial licence; location identifiers are already anonymised in SA-FARI. The negative result carries a responsible-deployment message: because label-free distance does *not* reliably flag when SAM3 will fail, conservation teams should not treat low-distance-from-training as a safety guarantee, and reliability should be checked with labelled spot-audits rather than assumed.

Chapter 6

Conclusion

6.1 Summary of findings

We asked whether SAM 3’s zero-shot promptable-tracking accuracy on an unseen (species, place) can be predicted, before running the model and without target labels, from four label-free distances — separately for detection and association. Under a deliberately powerful out-of-sample test the answer is no: no distance beats a mean-predictor baseline on either a species or a location hold-out, the one apparent effect is an animal-size confound, and a follow-up intrinsic-difficulty pivot collapses to the same size effect. Only object size predicts detection, modestly, and it is the nuisance the analysis controls for. Association carries almost no signal distinct from detection. The result is a rigorous, decomposed null.

6.2 Contributions

To our knowledge this is the first study to test, out of sample, whether a *promptable video tracker’s* zero-shot transfer is predictable from label-free, before-running signals, and the first to decompose that question into detection and association. Its contribution is the honest negative answer and its explanation: transfer is not distance-driven; the apparent signals reduce to an animal-size confound; and the test is shown to detect that one real signal when present, so the null reflects an absence of signal, not of power.

6.3 Future work

Two directions remain open. First, representational probing — whether the target species is separable in SAM 3’s own feature space — as a hedge against the unknown pretraining corpus. Second, broadening beyond a single dataset and backbone, and to targets with richer association structure (crowded, multi-object scenes) where the detection–association decomposition would have more to distinguish.

Bibliography

- [1] Christina Baek, Yiding Jiang, Aditi Raghunathan, and J. Zico Kolter. Agreement-on-the-line: Predicting the performance of neural networks under distribution shift. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2206.13089.
- [2] Sara Beery, Grant Van Horn, and Pietro Perona. Recognition in terra incognita. In *European Conference on Computer Vision (ECCV)*, 2018. arXiv:1807.04975.
- [3] Otto Brookes, Maksim Kukushkin, Majid Mirmehdi, et al. The PanAf-FGBG dataset: Understanding the impact of backgrounds in wildlife behaviour recognition. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. arXiv:2502.21201.
- [4] Nicolas Carion, Laura Gustafson, Yuan-Ting Hu, et al. SAM 3: Segment anything with concepts, 2025. arXiv:2511.16719.
- [5] Weijian Deng and Liang Zheng. Are labels always necessary for classifier accuracy evaluation? In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. arXiv:2007.02915.
- [6] Saurabh Garg, Sivaraman Balakrishnan, Zachary C. Lipton, Behnam Neyshabur, and Hanie Sedghi. Leveraging unlabeled data to predict out-of-distribution performance. In *International Conference on Learning Representations (ICLR)*, 2022. arXiv:2201.04234.
- [7] Sooyoung Jeon, Hongjie Tian, Lemeng Wang, Zheda Mai, Wei-Lun Chao, et al. Lessons and open questions from a unified study of camera-trap species recognition over time, 2026. arXiv:2603.20509.
- [8] Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, et al. WILDS: A benchmark of in-the-wild distribution shifts. In *International Conference on Machine Learning (ICML)*, 2021. arXiv:2012.07421.
- [9] Jonathon Luiten, Aljoša Ošep, Patrick Dendorfer, Philip Torr, Andreas Geiger, Laura Leal-Taixé, and Bastian Leibe. HOTA: A higher order metric for evaluating multi-object tracking. *International Journal of Computer Vision*, 129:548–578, 2021. arXiv:2009.07736.
- [10] John Miller, Rohan Taori, Aditi Raghunathan, Shiori Sagawa, Pang Wei Koh, Vaishaal Shankar, Percy Liang, Yair Carmon, and Ludwig Schmidt. Accuracy on the line: On the strong correlation between out-of-distribution and in-distribution generalization. In *International Conference on Machine Learning (ICML)*, 2021. arXiv:2107.04649.
- [11] Maxime Oquab, Timothée Darcet, Théo Moutakanni, et al. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research (TMLR)*, 2024. arXiv:2304.07193.
- [12] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning (ICML)*, 2021. arXiv:2103.00020.

- [13] Samuel Stevens, Jiaman Wu, Matthew J. Thompson, Elizabeth G. Campolongo, Chan Hee Song, David Edward Carlyn, et al. BioCLIP: A vision foundation model for the tree of life. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2311.18803.
- [14] Vanya V. Valindria, Ioannis Lavdas, Wenjia Bai, Konstantinos Kamnitsas, Eric O. Aboagye, Andrea G. Rockall, Daniel Rueckert, and Ben Glocker. Reverse classification accuracy: Predicting segmentation performance in the absence of ground truth. *IEEE Transactions on Medical Imaging*, 36(8):1597–1606, 2017. arXiv:1702.03407.
- [15] Dante Francisco Wasmuht, Otto Brookes, Maximillian Schall, et al. The SA-FARI dataset: Segment anything in footage of animals for recognition and identification. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. arXiv:2511.15622.
- [16] Seungju Yoo, Hyuk Kwon, Joong-Won Hwang, and Kibok Lee. Automated model evaluation for object detection via prediction consistency and reliability. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. arXiv:2508.12082.

Appendix A

Mathematical definitions

This appendix collects the exact definitions used in Chapter 3. The tracking metric is computed by the official evaluator and never re-implemented; the distances and the statistical procedures are ours and are stated here as they appear in the code.

A.1 The promptable-tracking metric

Promptable HOTA decomposes, at each localisation threshold α , into a detection and an association term [9]:

$$\text{HOTA}_\alpha = \sqrt{\text{DetA}_\alpha \cdot \text{AssA}_\alpha}, \quad \text{pHOTA} = \frac{1}{|\mathcal{A}|} \sum_{\alpha \in \mathcal{A}} \text{HOTA}_\alpha, \quad (\text{A.1})$$

with $\mathcal{A} = \{0.05, 0.10, \dots, 0.95\}$. We analyse the two components separately as the targets $\text{pDetA} = \text{DetA}$ (did it *find* the animals?) and $\text{pAssA} = \text{AssA}$ (did it *keep* each identity?), both bounded in $[0, 1]$. Scores are computed per cell by the vendored **VEval** evaluator.

A.2 Label-free distances

Let the frozen reference be the “seen” set of a split; every distance below is computed against that reference only.

Taxonomic distance. Write each species’ seven-level path $\pi(s) = (\text{kingdom}, \dots, \text{species})$, defined only for species with a complete taxonomy. Let $\text{shared}(a, b)$ be the length of the longest common *non-empty* prefix of two paths. The tree-step distance and the nearest-reference distance are

$$\tau(a, b) = L - \text{shared}(a, b) \quad (L = 7), \quad d_{\text{tax}}(s) = \min_{s' \in \mathcal{R}, s' \neq s} \tau(\pi(s), \pi(s')), \quad (\text{A.2})$$

where the exclusion $s' \neq s$ is the leave-one-species-out rule (Split A). Shared genus gives $\tau = 1$, family 2, order 3, fully disjoint 7; species without a full taxonomy take $d_{\text{tax}} = \text{NaN}$.

Prototype and nearest cosine distance. Given a set of ℓ_2 -normalised embeddings $\{v_i\}_{i=1}^N$, the prototype and the distance from a query v to a set of prototypes $\{p_j\}$ (optionally excluding an index e) are

$$p = \frac{\bar{v}}{\max(\|\bar{v}\|_2, 10^{-12})}, \quad \bar{v} = \frac{1}{N} \sum_i v_i, \quad d(v) = 1 - \max_{j \neq e} v^\top p_j. \quad (\text{A.3})$$

Because v and each p_j are unit-norm, $v^\top p_j$ is the cosine similarity, so $d \in [0, 2]$.

Visual and environment distance. Both use DINOv2 embeddings. The *visual* distance embeds the ground-truth-mask-cropped animal (the background is zeroed before cropping), forms one prototype per species, and sets $d_{\text{vis}}(s)$ to the nearest *other* species’ prototype. The *environment* distance embeds the animal-masked-out background (the animal union is replaced by the mean colour of the remaining pixels), forms one prototype per location, and sets $d_{\text{env}}(\ell)$ to the nearest other location. Two interpretable scene covariates accompany the environment distance:

$$\text{achromatic}(f) = \left[\overline{\max_c f_c - \min_c f_c} < \tau_{\text{ir}} \right], \quad \tau_{\text{ir}} = 12, \quad (\text{A.4})$$

$$\text{is_night_ir}(\ell) = \left[\overline{\text{achromatic}(f)} > 0.5 \right], \quad \text{clutter}(\ell) = \overline{\text{Var}(\nabla^2 f)}, \quad (\text{A.5})$$

where the bar denotes a mean over a location’s frames, $\max_c - \min_c$ is the per-pixel channel spread, and ∇^2 is the four-neighbour Laplacian on the greyscale background.

Temporal gap and animal size. For a cell with footage year y and reference years $\mathcal{Y}$, and for the size covariate,

$$d_{\text{temp}} = \min_{r \in \mathcal{Y}} |y - r|, \quad \log_area(s) = \ln(1 + \overline{area(s)}), \quad (\text{A.6})$$

where $\overline{area(s)}$ is the mean ground-truth mask pixel-count over sampled masklets of species s . Size is a *nuisance* covariate, not a hypothesis variable.

A.3 Model and statistics

Support-weighted fractional-logit GLM. Each bounded target $y \in [0, 1]$ is modelled with a Binomial (logit-link) GLM

$$\text{logit } \mathbb{E}[y] = x^\top \beta, \quad \text{fit with variance weights } w = n_{\text{frames}}, \quad (\text{A.7})$$

where the design x standardises each continuous predictor $z = (x - \mu)/\sigma$ (missing values imputed to the training mean μ , zero-variance columns dropped), adds $\log_{\text{support}}$, the standardised $\ln(1 + n_{\text{frames}})$, passes binary `is_night_ir` through, and adds an intercept. Standardisation statistics are learned on training rows only.

Explained deviance and the cross-validation gain. Fit quality is the McFadden pseudo- R^2 , and per-cell out-of-sample gain over a mean predictor is

$$R_{\text{McF}}^2 = 1 - \frac{D}{D_0}, \quad \delta_i = \underbrace{|y_i - \bar{y}|}_{\text{baseline error}} - \underbrace{|y_i - \hat{y}_i|}_{\text{model error}}, \quad (\text{A.8})$$

with D, D_0 the model and null (intercept-only) deviances; $\delta_i > 0$ means the distance model beats “guess the mean” on cell i .

Weighted Pearson correlation. For the species-level false-positive analysis, with weights w (per-species probe counts),

$$r_w = \frac{\sum_i w_i (x_i - \bar{x}_w)(y_i - \bar{y}_w)}{\sqrt{\sum_i w_i (x_i - \bar{x}_w)^2} \sqrt{\sum_i w_i (y_i - \bar{y}_w)^2}}, \quad \bar{x}_w = \frac{\sum_i w_i x_i}{\sum_i w_i}. \quad (\text{A.9})$$

Appendix B

Algorithms

Pseudocode for the four non-obvious procedures. The two bootstraps differ deliberately: the cross-validation gain (Alg. 3) is tested *one-sided* (does the model beat the mean?), whereas the false-positive correlation (Alg. 4) is tested *two-sided* (does hallucination change with novelty, in either direction?).

Algorithm 1: Memory-lean, cache-first crop embedding (`embed_crops`).

Input: items (key, record, frame, crop_fn); embedding cache; embedder; chunk size C , workers W

Output: $\{\text{key} \mapsto \text{embedding}\}$

```
1 misses ← [item : cache.get(item.key) = ∅];
2 for each contiguous chunk  $B$  of  $C$  items in misses do
3   group  $B$  by video; for each video do ensure_frames(missing frames);
4   loaded ← map load_frame → crop_fn over  $B$  using  $W$  threads;
5   pending ← [( $k, c$ ) ∈ loaded :  $c \neq \emptyset$ ];
6   if pending ≠ ∅ then
7     vecs ← embedder.embed(crops of pending);
8     for ( $k, \_$ ),  $v$  in pending, vecs do cache.put( $k, v$ );
9   free pending ;                                // bound peak RAM to one chunk
10  every 16 chunks: cache.save() ;                 // resumable checkpoint
11 if misses ≠ ∅ then cache.save();
12 return {item.key ↦ cache.get(item.key)} for all items;
```

Algorithm 2: Group-cluster-bootstrap coefficient CIs (`group_bootstrap_cis`).

Input: scored cells; target; grouping columns $\mathcal{G} = (\text{category_id}, \text{location_id})$; resamples B ; level $\alpha = 0.05$
Output: per-coefficient interval $[\text{ci_lo}, \text{ci_hi}]$

- 1 **foreach** grouping column $g \in \mathcal{G}$ present in the data **do**
- 2 $U \leftarrow$ unique groups of g ; **if** $|U| < 2$ **then** skip;
- 3 **for** $b \leftarrow 1$ to B **do**
- 4 draw $|U|$ groups from U with replacement; sample $\leftarrow$ all rows of the drawn groups;
- 5 refit the GLM on sample; record its coefficient vector ; // skip a singular resample
- 6 $\log \leftarrow$ percentile $\frac{\alpha}{2}$, $\text{hi}_g \leftarrow$ percentile $1 - \frac{\alpha}{2}$, per coefficient;
- 7 **return** conservative envelope $\text{ci_lo} = \min_g \log$, $\text{ci_hi} = \max_g \text{hi}_g$;

Algorithm 3: Paired group-bootstrap cross-validation significance (`oos_predictions + _bootstrap_delta`).

Input: scored cells; target; group column (species or location); resamples B
Output: ΔMAE , its 95% CI, and one-sided p

/* Leave-one-group-out out-of-sample prediction */

- 1 **foreach** held-out group **do**
- 2 refit the GLM on the other groups (own standardisation); predict the held-out cells ;
 // assert train/test groups disjoint
- 3 $\delta_i \leftarrow |y_i - \bar{y}| - |y_i - \hat{y}_i|$ for every predicted cell;
 /* Group bootstrap of the mean gain */
- 4 **for** $b \leftarrow 1$ to B **do**
- 5 draw whole groups with replacement; $\text{boot}_b \leftarrow \text{mean}(\delta \text{ over the drawn cells})$;
- 6 $\text{CI} \leftarrow$ percentiles $[2.5, 97.5]$ of boot ; $p \leftarrow \Pr(\text{boot} \leq 0)$; significant $\Leftrightarrow \text{CI_lo} > 0$;
- 7 **return** $\Delta\text{MAE} = \text{mean}(\delta)$, CI, p ;

Algorithm 4: False-positive / hallucination analysis (`false_positives`).

Input: split; score threshold t ; resamples B
Output: overall FP rate; per-species FP rate vs. novelty with a bootstrap CI

- 1 $\mathcal{H} \leftarrow$ hard-negative probes (queried species absent);
- 2 build a partition whose probe species are the hard-negative species; compute $d_{\text{tax}}, d_{\text{vis}}$ for them;
- 3 **foreach** probe (*video, species*) $\in \mathcal{H}$ with a prediction file **do**
- 4 $\text{fp} \leftarrow [\max_k \text{score}_k \geq t]$; // a masklet returned for an absent species = a hallucination
- 5 attach $d_{\text{tax}}, d_{\text{vis}}$ of the species;
- 6 aggregate to per-species FP rate; $r \leftarrow \text{weighted-Pearson}(\text{FP rate}, \text{distance})$;
- 7 **for** $b \leftarrow 1$ to B **do**
- 8 resample species with replacement; recompute r ;
- 9 $\text{CI} \leftarrow$ percentiles $[2.5, 97.5]$; significant $\Leftrightarrow \text{CI}$ excludes 0 (two-sided);
- 10 **return** overall FP rate, taxonomic-tercile breakdown, r and CI for taxonomic and visual novelty;

Appendix C

Supplementary results